

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

Class: Final Year (Computer Science and Engineering)

Year: 2022-23 **Semester:** 1

Course: High Performance Computing Lab

Practical No. 10

Exam Seat No: 2019BTECS00064

Name – Kunal Santosh Kadam

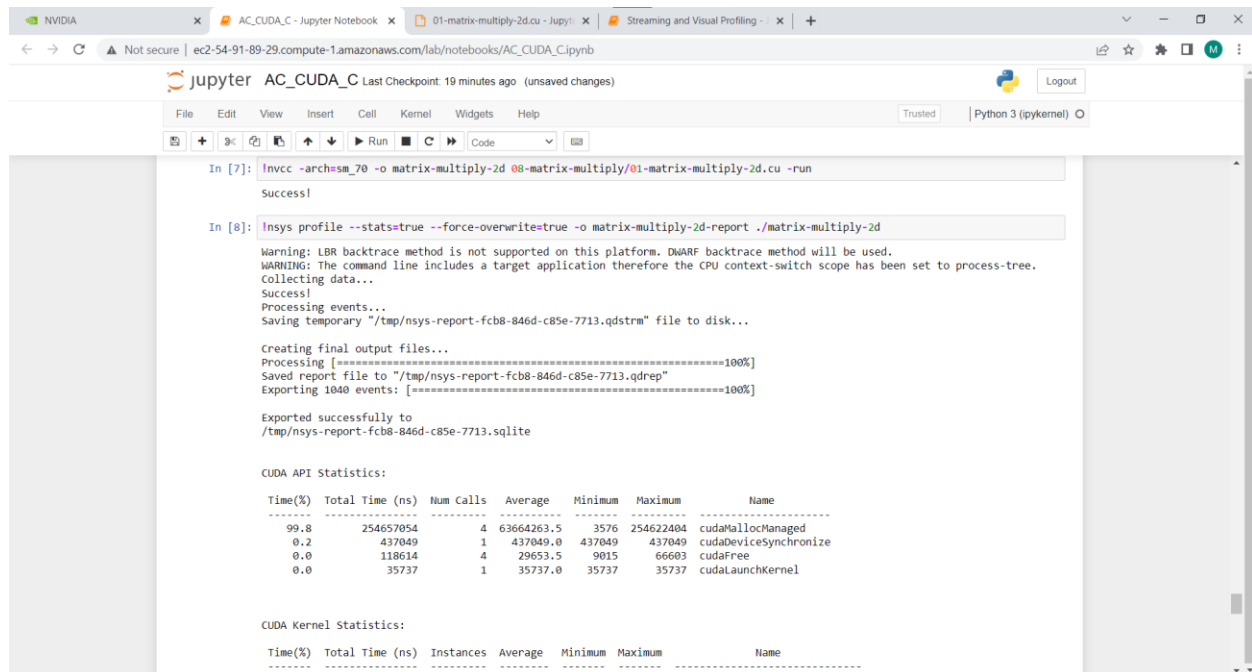
Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

Problem Statement 1:

Implement Matrix-matrix Multiplication using global memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute

Screenshot #:



```
In [7]: nvcc -arch=sm_70 -o matrix-multiply-2d 01-matrix-multiply/01-matrix-multiply-2d.cu -run
Success!

In [8]: nsys profile --stats=true --force-override=true -o matrix-multiply-2d-report ./matrix-multiply-2d
Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.
collecting data...
Success!
Processing events...
Saving temporary "/tmp/nsys-report-fcb8-846d-c85e-7713.qdstrm" file to disk...

Creating final output files...
Processing [=====100%]
saved report file to "/tmp/nsys-report-fcb8-846d-c85e-7713.qdrep"
Exporting 1040 events: [=====100%]

Exported successfully to
/tmp/nsys-report-fcb8-846d-c85e-7713.sqlite

CUDA API Statistics:
Time(%) Total Time (ns) Num Calls Average Minimum Maximum Name
-----
99.8 254657054 4 63664263.5 3576 254622404 cudaMallocManaged
0.2 437049 1 437049.0 437049 437049 cudaDeviceSynchronize
0.0 118614 4 29653.5 9015 66603 cudaFree
0.0 35737 1 35737.0 35737 35737 cudaLaunchKernel

CUDA Kernel Statistics:
Time(%) Total Time (ns) Instances Average Minimum Maximum Name
-----
```

Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

The screenshot displays a Jupyter Notebook interface with the following content:

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	434455	1	434455.0	434455	434455	matrixMulGPU(int*, int*, int*)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
53.9	13598	2	6799.0	4895	8703	[CUDA Unified Memory memcpy HtoD]
46.1	11615	2	5807.5	1471	10144	[CUDA Unified Memory memcpy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
64.000	2	32.000	4.000	60.000	[CUDA Unified Memory memcpy DtoH]
64.000	2	32.000	20.000	44.000	[CUDA Unified Memory memcpy HtoD]

Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
62.1	230567541	16	14410471.3	29318	100130911	poll
26.3	97752443	670	145899.2	1007	17139283	ioctl
10.2	37936772	14	2709769.4	11515	20502483	sem_timedwait
0.7	2651839	92	28824.3	1469	754433	mmap
0.5	1779608	82	21702.5	6245	35180	open64
0.1	215667	3	71889.0	69900	75844	fgets
0.0	160484	4	40121.0	33258	48363	pthread_create

Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
62.1	230567541	16	14410471.3	29318	100130911	poll
26.3	97752443	670	145899.2	1007	17139283	ioctl
10.2	37936772	14	2709769.4	11515	20502483	sem_timedwait
0.7	2651839	92	28824.3	1469	754433	mmap
0.5	1779608	82	21702.5	6245	35180	open64
0.1	215667	3	71889.0	69900	75844	fgets
0.0	160484	4	40121.0	33258	48363	pthread_create
0.0	101783	23	4425.3	1485	21643	fopen
0.0	82757	11	7523.4	4326	12111	write
0.0	31240	5	6248.0	3178	9558	open
0.0	30068	7	4295.4	1021	9700	fgetc
0.0	28901	7	4128.7	2435	7479	munmap
0.0	22450	16	1403.1	1049	2218	fclose
0.0	20880	13	1606.2	1034	3495	read
0.0	15274	2	7637.0	6477	8797	socket
0.0	11227	1	11227.0	11227	11227	sem_wait
0.0	9309	4	2327.3	1906	2792	mprotect
0.0	8897	1	8897.0	8897	8897	pipe2
0.0	7564	1	7564.0	7564	7564	connect
0.0	6936	6	1156.0	1057	1307	fcntl
0.0	2632	1	2632.0	2632	2632	bind
0.0	1566	1	1566.0	1566	1566	listen

Report file moved to "/dli/task/matrix-multiply-2d-report.qdrep"
Report file moved to "/dli/task/matrix-multiply-2d-report.sqlite"

```
In [2]: %js
var port = (window.location.port == 80) ? "" : (":" + window.location.port);
var url = 'http://' + window.location.hostname + port + '/nsight/vnc.html?resize=scale';
let a = document.createElement('a');
a.setAttribute('href', url);
a.setAttribute('target', '_blank');
```

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

Information #:

```
#include <stdio.h>

#define N 64

__global__ void matrixMulGPU( int * a, int * b, int * c )
{
    int val = 0;

    int row = blockIdx.x * blockDim.x + threadIdx.x;
    int col = blockIdx.y * blockDim.y + threadIdx.y;

    if (row < N && col < N)
    {
        for ( int k = 0; k < N; ++k )
            val += a[row * N + k] * b[k * N + col];
        c[row * N + col] = val;
    }
}

void matrixMulCPU( int * a, int * b, int * c )
{
    int val = 0;

    for( int row = 0; row < N; ++row )
        for( int col = 0; col < N; ++col )
        {
            val = 0;
            for ( int k = 0; k < N; ++k )
                val += a[row * N + k] * b[k * N + col];
            c[row * N + col] = val;
        }
}
```

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

```
int main()
{
    int *a, *b, *c_cpu, *c_gpu;

    int size = N * N * sizeof (int); // Number of bytes of an N x N matrix

    // Allocate memory
    cudaMallocManaged (&a, size);
    cudaMallocManaged (&b, size);
    cudaMallocManaged (&c_cpu, size);
    cudaMallocManaged (&c_gpu, size);

    // Initialize memory
    for( int row = 0; row < N; ++row )
        for( int col = 0; col < N; ++col )
        {
            a[row*N + col] = row;
            b[row*N + col] = col+2;
            c_cpu[row*N + col] = 0;
            c_gpu[row*N + col] = 0;
        }

    dim3 threads_per_block (16, 16, 1); // A 16 x 16 block threads
    dim3 number_of_blocks ((N / threads_per_block.x) + 1, (N /
threads_per_block.y) + 1, 1);

    matrixMulGPU <<< number_of_blocks, threads_per_block >>> ( a, b,
c_gpu );

    cudaDeviceSynchronize(); // Wait for the GPU to finish before
proceeding

    // Call the CPU version to check our work
```

```
matrixMulCPU( a, b, c_cpu );
```

```
// Compare the two answers to make sure they are equal
```

```
bool error = false;
```

```
for( int row = 0; row < N && !error; ++row )
```

```
    for( int col = 0; col < N && !error; ++col )
```

```
        if (c_cpu[row * N + col] != c_gpu[row * N + col])
```

```
        {
```

```
            printf("FOUND ERROR at c[%d][%d]\n", row, col);
```

```
            error = true;
```

```
            break;
```

```
        }
```

```
if (!error)
```

```
    printf("Success!\n");
```

```
// Free all our allocated memory
```

```
cudaFree(a);
```

```
cudaFree(b);
```

```
cudaFree( c_cpu );
```

```
cudaFree( c_gpu );
```

```
}
```

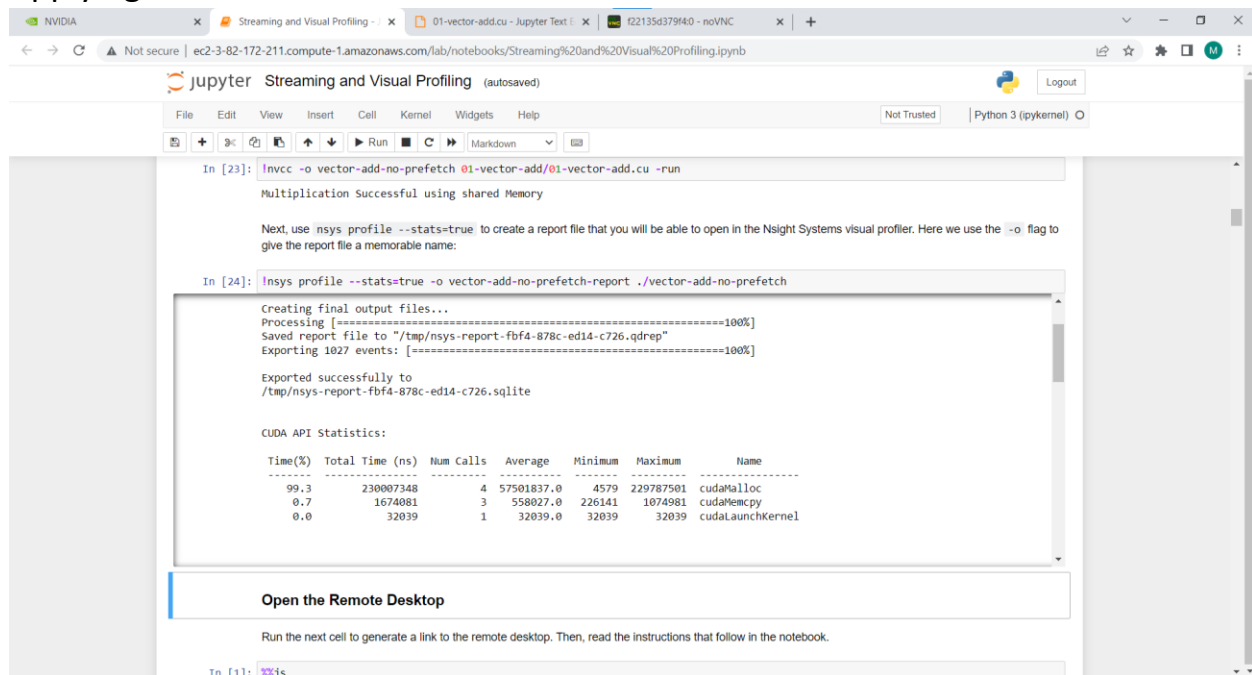
Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

Problem Statement 2:

Implement Matrix-Matrix Multiplication using shared memory in CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

Screenshot #:

Applying 512 blocks with 32 threads each block



The screenshot shows a Jupyter Notebook interface with the following content:

```
In [23]: nvcc -o vector-add-no-prefetch 01-vector-add/01-vector-add.cu -run
```

Multiplication Successful using shared Memory

Next, use `nsys profile --stats=true` to create a report file that you will be able to open in the Nsight Systems visual profiler. Here we use the `-o` flag to give the report file a memorable name:

```
In [24]: nsys profile --stats=true -o vector-add-no-prefetch-report ./vector-add-no-prefetch
```

Creating final output files...
Processing [=====100%]
Saved report file to "/tmp/nsys-report-fbf4-878c-ed14-c726.qdrep"
Exporting 1027 events: [=====100%]
Exported successfully to
/tmp/nsys-report-fbf4-878c-ed14-c726.sqlite

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.3	230007348	4	57501837.0	4579	229787501	cudaMalloc
0.7	1674081	3	558027.0	226141	1074981	cudaMemcpy
0.0	32039	1	32039.0	32039	32039	cudaLaunchKernel

Open the Remote Desktop

Run the next cell to generate a link to the remote desktop. Then, read the instructions that follow in the notebook.

```
In [1]: %x
```

Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

The image displays a Jupyter Notebook interface for 'Streaming and Visual Profiling' (autosaved) running on a remote desktop. The notebook shows the execution of two code cells. The first cell runs `nvcc -o vector-add-no-prefetch 01-vector-add/01-vector-add.cu -run`, resulting in 'Multiplication Successful using shared Memory'. The second cell runs `nsys profile --stats=true -o vector-add-no-prefetch-report ./vector-add-no-prefetch`, which generates performance statistics.

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	209471	1	209471.0	209471	209471	MatrixMulSh(float*, float*, float*, int)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
68.0	323870	2	161935.0	163727	162143	[CUDA memcopy HtoD]
32.0	152734	1	152734.0	152734	152734	[CUDA memcopy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Time(%)	Total Size (KiB)	Operations	Average	Minimum	Maximum	Operation
68.0	16384	2	8192.0	8192	8192	[CUDA memcopy HtoD]
32.0	16384	1	16384.0	16384	16384	[CUDA memcopy DtoH]

Open the Remote Desktop

Run the next cell to generate a link to the remote desktop. Then, read the instructions that follow in the notebook.

The second screenshot shows the NVIDIA Systems 2021.3.1 interface. The 'Timeline View' displays a horizontal timeline with various colored bars representing different system components and their execution times. The 'Events View' is currently empty, with a note indicating that right-clicking a timeline row and selecting 'Show in Events View' will populate it.

Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

Applying 256 blocks with 16 threads each block

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.3	230319260	4	57579815.0	4423	230097613	cudaMalloc
0.7	1600633	3	533544.3	247424	980856	cudaMemcpy
0.0	24192	1	24192.0	24192	24192	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	144799	1	144799.0	144799	144799	MatrixMulSh(float*, float*, float*, int)

Open the Remote Desktop

Run the next cell to generate a link to the remote desktop. Then, read the instructions that follow in the notebook.

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
67.9	323742	2	161871.0	161759	161983	[CUDA memcpy HtoD]
32.1	152703	1	152703.0	152703	152703	[CUDA memcpy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
976.563	1	976.563	976.563	976.563	[CUDA memcpy DtoH]
1953.125	2	976.563	976.563	976.563	[CUDA memcpy HtoD]

Operating System Runtime API Statistics:

Open the Remote Desktop

Run the next cell to generate a link to the remote desktop. Then, read the instructions that follow in the notebook.

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

Information #:

```
#include <stdio.h>
#include <math.h>
#define TILE_WIDTH 2

/*matrix multiplication kernels*/

// shared
__global__ void
MatrixMulSh( float *Md , float *Nd , float *Pd , const int WIDTH )
{

    //Taking shared array to break the MAtrix in Tile widht and fatch
    them in that array per ele

    __shared__ float Mds [TILE_WIDTH][TILE_WIDTH] ;

    __shared__ float Nds [TILE_WIDTH][TILE_WIDTH] ;

    // calculate thread id
    unsigned int col = TILE_WIDTH*blockIdx.x + threadIdx.x ;
    unsigned int row = TILE_WIDTH*blockIdx.y + threadIdx.y ;

    for (int m = 0 ; m<WIDTH/TILE_WIDTH ; m++ ) // m indicate number
    of phase
    {
        Mds[threadIdx.y][threadIdx.x] = Md[row*WIDTH +
        (m*TILE_WIDTH + threadIdx.x)] ;
        Nds[threadIdx.y][threadIdx.x] = Nd[ ( m*TILE_WIDTH +
        threadIdx.y ) * WIDTH + col] ;
        __syncthreads() ; // for synchronizeing the threads

        // Do for tile
    }
}
```

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

```
        for ( int k = 0; k<TILE_WIDTH ; k++ )
            Pd[row*WIDTH + col]+= Mds[threadIdx.x][k] *
Nds[k][threadIdx.y] ;
        __syncthreads() ; // for synchronizeing the threads

    }
}

// main routine
int main ()
{
    const int WIDTH = 500;
    float array1_h[WIDTH][WIDTH] ,array2_h[WIDTH][WIDTH],
M_result_array_h[WIDTH][WIDTH] ;
    float *array1_d , *array2_d ,*result_array_d ,*M_result_array_d ; //
device array
    int i , j ;
    //input in host array
    for ( i = 0 ; i<WIDTH ; i++ )
    {
        for (j = 0 ; j<WIDTH ; j++ )
        {
            array1_h[i][j] = (i + 2*j) %500 ;
            array2_h[i][j] = (i + 3*j) %500 ;
        }
    }

    //create device array cudaMalloc ( (void **)&array_name,
sizeofmatrixinbytes) ;

    cudaMalloc((void **) &array1_d , WIDTH*WIDTH*sizeof (int) ) ;

    cudaMalloc((void **) &array2_d , WIDTH*WIDTH*sizeof (int) ) ;
```

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

```
//copy host array to device array; cudaMemcpy ( dest , source ,  
WIDTH , direction )  
  
    cudaMemcpy ( array1_d , array1_h , WIDTH*WIDTH*sizeof (int) ,  
cudaMemcpyHostToDevice ) ;  
  
    cudaMemcpy ( array2_d , array2_h , WIDTH*WIDTH*sizeof (int) ,  
cudaMemcpyHostToDevice ) ;  
  
//allocating memory for resultant device array  
  
    cudaMalloc((void **) &result_array_d , WIDTH*WIDTH*sizeof (int) ) ;  
  
    cudaMalloc((void **) &M_result_array_d , WIDTH*WIDTH*sizeof  
(int) ) ;  
  
    MatrixMulSh<<<512,32>>> ( array1_d , array2_d ,M_result_array_d ,  
WIDTH) ;  
  
    // all gpu function blocked till kernel is working  
    //copy back result_array_d to result_array_h  
  
    cudaMemcpy(M_result_array_h , M_result_array_d ,  
WIDTH*WIDTH*sizeof(int) ,cudaMemcpyDeviceToHost) ;  
  
    printf("Multiplication Successful using shared Memory");  
  
}
```

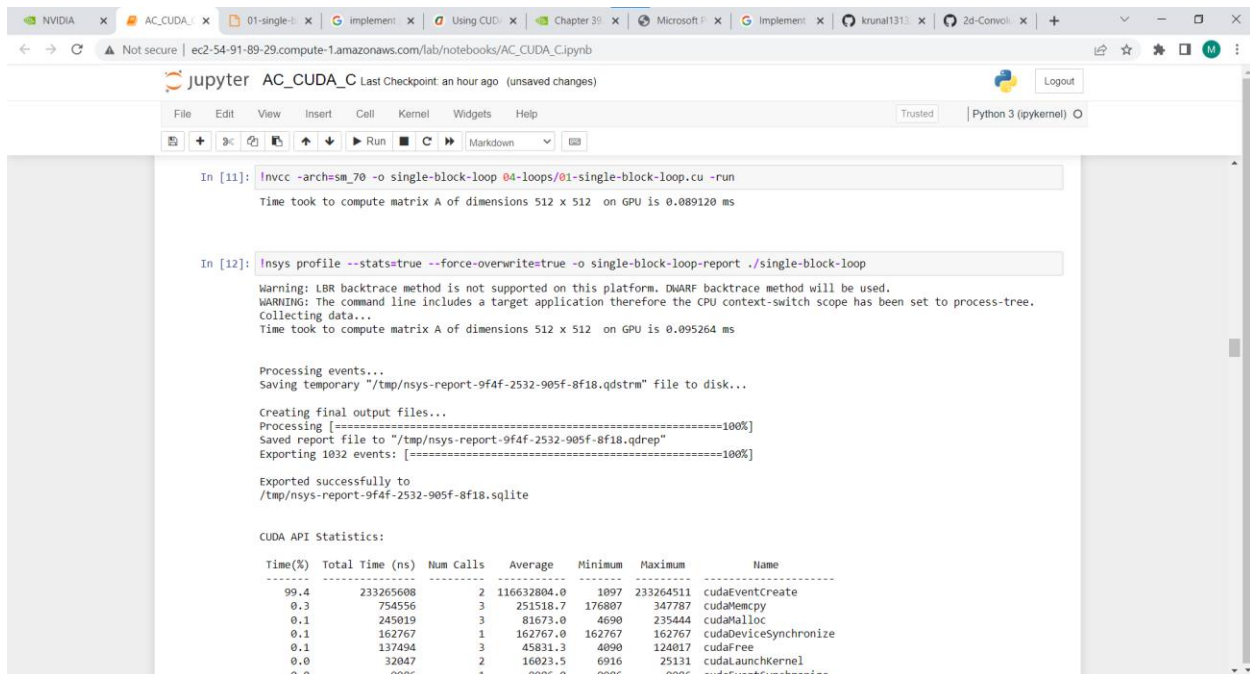
Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

Problem Statement 4:

Implement 2D Convolution using shared memory using CUDA C. Analyze and tune the program for getting maximum speed up. Do Profiling and state what part of the code takes the huge amount of time to execute.

Screenshot #:



The screenshot shows a Jupyter Notebook interface with the following content:

```
In [11]: nvcc -arch=sm_70 -o single-block-loop 04-loops/01-single-block-loop.cu -run
Time took to compute matrix A of dimensions 512 x 512 on GPU is 0.089120 ms
```

```
In [12]: nsys profile --stats=true --force-override=true -o single-block-loop-report ./single-block-loop

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.
Collecting data...
Time took to compute matrix A of dimensions 512 x 512 on GPU is 0.095264 ms

Processing events...
Saving temporary "/tmp/nsys-report-9f4f-2532-905f-8f18.qdstrm" file to disk...

Creating final output files...
Processing [=====100%]
Saved report file to "/tmp/nsys-report-9f4f-2532-905f-8f18.qdrep"
Exporting 1032 events: [=====100%]

Exported successfully to
/tmp/nsys-report-9f4f-2532-905f-8f18.sqlite

CUDA API Statistics:
```

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.4	233265608	2	116632804.0	1097	233264511	cudaEventCreate
0.3	754556	3	251518.7	176887	347787	cudaMemcpy
0.1	245019	3	81673.0	4690	235444	cudaMalloc
0.1	162767	1	162767.0	162767	162767	cudaDeviceSynchronize
0.1	137494	3	45831.3	4090	124017	cudaFree
0.0	32047	2	16023.5	6916	25131	cudaLaunchKernel
0.0	9086	1	9086.0	9086	9086	cudaEventSynchronize

Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

```
Exported successfully to
/tmp/nsys-report-9f4f-2532-905f-8f18.sqlite

CUDA API Statistics:

Time(%) Total Time (ns) Num Calls Average Minimum Maximum Name
99.4 233265608 2 116632804.0 1097 233264511 cudaEventCreate
0.3 754556 3 251518.7 176807 347787 cudaMemcpy
0.1 245019 3 81673.0 4690 235444 cudaMalloc
0.1 162767 1 162767.0 162767 162767 cudaDeviceSynchronize
0.1 137494 3 45831.3 4090 124017 cudaFree
0.0 32047 2 16023.5 6916 25131 cudaLaunchKernel
0.0 9986 1 9986.0 9986 9986 cudaEventSynchronize
0.0 7682 2 3841.0 2286 5476 cudaEventRecord

CUDA Kernel Statistics:

Time(%) Total Time (ns) Instances Average Minimum Maximum Name
100.0 171900 2 85950.0 85950 85950 Convolution(float*, float*, float*, int, int, int, int, i
nt)

CUDA Memory Operation Statistics (by time):

Time(%) Total Time (ns) Operations Average Minimum Maximum Operation
51.5 170813 2 85406.5 1056 169757 [CUDA memcpy HtoD]
48.5 160668 1 160668.0 160668 160668 [CUDA memcpy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Total Operations Average Minimum Maximum Operation
1024.035 2 512.018 0.035 1024.000 [CUDA memcpy HtoD]
1016.016 1 1016.016 1016.016 1016.016 [CUDA memcpy DtoH]

Operating System Runtime API Statistics:

Time(%) Total Time (ns) Num Calls Average Minimum Maximum Name
64.1 200708335 13 15439102.7 33134 98135208 poll
34.1 106857144 665 160687.4 1003 19107269 ioctl
0.8 2640025 89 29663.2 1496 761498 mmap
0.5 1501697 82 18313.4 7415 30536 open64
0.1 426595 10 42659.5 23273 123152 sem_timedwait
0.1 229100 3 76366.7 71081 80950 fgetc
0.1 162914 4 40728.5 31774 48271 pthread_create
0.0 119135 23 5179.8 1673 27930 fopen
0.0 88112 11 8010.2 4201 14533 write
0.0 37126 5 7425.2 5436 9432 open
0.0 29922 13 2301.7 1395 4105 read
0.0 25908 16 1619.3 1151 3329 fclose
0.0 25890 6 4315.0 1040 12105 fgetc
0.0 23854 6 3975.7 1946 5779 munmap
0.0 15782 2 7891.0 6216 9566 socket
0.0 9021 1 9021.0 9021 9021 pipe2
0.0 8641 4 2160.3 1779 2744 mprotect
0.0 8608 1 8608.0 8608 8608 connect
0.0 8596 7 1228.0 1082 1374fcntl
0.0 2372 1 2372.0 2372 2372 bind
0.0 2192 1 2192.0 2192 2192 sem_wait
0.0 1626 1 1626.0 1626 1626 listen

Report file moved to "/dli/task/single-block-loop-report.qdrep"
Report file moved to "/dli/task/single-block-loop-report.sqlite"
```

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

Information #:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define BLOCK_SIZE 32
#define WA 512
#define HA 512
#define HC 3
#define WC 3
#define WB (WA - WC + 1)
#define HB (HA - HC + 1)

__global__ void Convolution(float* A, float* B, float* C, int numRows, int
numACols, int numBRows, int numBCols, int numCRows, int numCCols)
{
    int col = blockIdx.x * (BLOCK_SIZE - WC + 1) + threadIdx.x;
    int row = blockIdx.y * (BLOCK_SIZE - WC + 1) + threadIdx.y;
    int row_i = row - WC + 1;
    int col_i = col - WC + 1;

    float tmp = 0;

    __shared__ float shm[BLOCK_SIZE][BLOCK_SIZE];

    if (row_i < WA && row_i >= 0 && col_i < WA && col_i >= 0)
    {
        shm[threadIdx.y][threadIdx.x] = A[col_i * WA + row_i];
    }
    else
    {
        shm[threadIdx.y][threadIdx.x] = 0;
    }
}
```

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

```
}

__syncthreads();

if (threadIdx.y < (BLOCK_SIZE - WC + 1) && threadIdx.x < (BLOCK_SIZE
- WC + 1) && row < (WB - WC + 1) && col < (WB - WC + 1))
{
    for (int i = 0; i < WC; i++)
        for (int j = 0; j < WC; j++)
            tmp += shm[threadIdx.y + i][threadIdx.x + j] *
C[j*WC + i];
    B[col*WB + row] = tmp;
}
}

void randomInit(float* data, int size)
{
    for (int i = 0; i < size; ++i)
        data[i] = rand() / (float)RAND_MAX;
}

int main(int argc, char** argv)
{
    srand(2006);
    cudaError_t error;
    cudaEvent_t start_G, stop_G;

    cudaEventCreate(&start_G);
    cudaEventCreate(&stop_G);

    unsigned int size_A = WA * HA;
    unsigned int mem_size_A = sizeof(float) * size_A;
    float* h_A = (float*)malloc(mem_size_A);
```


Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

```
unsigned int size_B = WB * HB;
unsigned int mem_size_B = sizeof(float) * size_B;
float* h_B = (float*)malloc(mem_size_B);

unsigned int size_C = WC * HC;
unsigned int mem_size_C = sizeof(float) * size_C;
float* h_C = (float*)malloc(mem_size_C);

randomInit(h_A, size_A);
randomInit(h_C, size_C);

float* d_A;
float* d_B;
float* d_C;

error = cudaMalloc((void**)&d_A, mem_size_A);
if (error != cudaSuccess)
{
    fprintf(stderr, "GPUassert: %s in cudaMalloc for A\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}

error = cudaMalloc((void**)&d_B, mem_size_B);
if (error != cudaSuccess)
{
    fprintf(stderr, "GPUassert: %s in cudaMalloc for B\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}

error = cudaMalloc((void**)&d_C, mem_size_C);
if (error != cudaSuccess)
```

Walchand College of Engineering, Sangli
Department of Computer Science and Engineering

```
{
    fprintf(stderr, "GPUassert: %s in cudaMalloc for C\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}

error = cudaMemcpy(d_A, h_A, mem_size_A,
cudaMemcpyHostToDevice);
if (error != cudaSuccess)
{
    fprintf(stderr, "GPUassert: %s in cudaMemcpy for A\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}

error = cudaMemcpy(d_C, h_C, mem_size_C,
cudaMemcpyHostToDevice);
if (error != cudaSuccess)
{
    fprintf(stderr, "GPUassert: %s in cudaMemcpy for C\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}

dim3 threads(BLOCK_SIZE, BLOCK_SIZE);
dim3 grid((WB - 1) / (BLOCK_SIZE - WC + 1), (WB - 1) / (BLOCK_SIZE -
WC + 1));

Convolution << < grid, threads >> >(d_A, d_B, d_C, HA, WA, HB, WB,
HC, WC);

cudaEventRecord(start_G);
```

Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

```
Convolution << < grid, threads >> >(d_A, d_B, d_C, HA, WA, HB, WB,
HC, WC);
error = cudaGetLastError();
if (error != cudaSuccess)
{
    fprintf(stderr, "GPUassert: %s in launching kernel\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}

error = cudaDeviceSynchronize();

if (error != cudaSuccess)
{
    fprintf(stderr, "GPUassert: %s in cudaDeviceSynchronize\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}

cudaEventRecord(stop_G);

cudaEventSynchronize(stop_G);

error = cudaMemcpy(h_B, d_B, mem_size_B,
cudaMemcpyDeviceToHost);

if (error != cudaSuccess)
{
    fprintf(stderr, "GPUassert: %s in cudaMemcpy for B\n",
cudaGetErrorString(error));
    return EXIT_FAILURE;
}
```

Walchand College of Engineering, Sangli

Department of Computer Science and Engineering

```
float milliseconds = 0;
```

```
cudaEventElapsedTime(&milliseconds, start_G, stop_G);
```

```
printf("Time took to compute matrix A of dimensions %d x %d on  
GPU is %f ms \n \n \n", WA, HA, milliseconds);
```

```
//for (int i = 0;i < HB;i++)
```

```
//{
```

```
//    for (int j = 0;j < WB;j++)
```

```
//    {
```

```
//        printf("%f ", h_B[i*HB + j]);
```

```
//    }
```

```
//    printf("\n");
```

```
//}
```

```
free(h_A);
```

```
free(h_B);
```

```
free(h_C);
```

```
cudaFree(d_A);
```

```
cudaFree(d_B);
```

```
cudaFree(d_C);
```

```
return EXIT_SUCCESS;
```

```
}
```

Github Link:

<https://github.com/Kunalkadam179/HPC-Assignment/tree/main/Assignment%20-%2010>